



Open in app

Get started



Published in Towards Data Science



Susan Li

Follow

Sep 24, 2017 · 6 min read · Listen



Save



# A Gentle Introduction on Market Basket Analysis — Association Rules

$$\begin{aligned}
 \text{Rule: } X \Rightarrow Y & \begin{cases} \text{Support} = \frac{\text{freq}(X, Y)}{N} \\ \text{Confidence} = \frac{\text{freq}(X, Y)}{\text{freq}(X)} \\ \text{Lift} = \frac{\text{Support}}{\text{Supp}(X) \times \text{Supp}(Y)} \end{cases}
 \end{aligned}$$

*Example:*

Rule	Support	Confidence	Lift
$A \Rightarrow D$	2/5	2/3	10/9
$C \Rightarrow A$	2/5	2/4	5/6
$A \Rightarrow C$	2/5	2/3	5/6
$B \& C \Rightarrow D$	1/5	1/3	5/9

Source: UofT

## Introduction

Market Basket Analysis is one of the key techniques used by large retailers to uncover associations between items. It works by looking for combinations of items that occur together frequently in transactions. To put it another way, it allows retailers to identify



[Open in app](#)[Get started](#)

interestingness, based on the concept of strong rules.

## An example of Association Rules

- Assume there are 100 customers
- 10 of them bought milk, 8 bought butter and 6 bought both of them.
- bought milk => bought butter
- support =  $P(\text{Milk \& Butter}) = 6/100 = 0.06$
- confidence =  $\text{support}/P(\text{Butter}) = 0.06/0.08 = 0.75$
- lift =  $\text{confidence}/P(\text{Milk}) = 0.75/0.10 = 7.5$

Note: this example is extremely small. In practice, a rule needs the support of several hundred transactions, before it can be considered statistically significant, and datasets often contain thousands or millions of transactions.

Ok, enough for the theory, let's get to the code.

The dataset we are using today comes from [UCI Machine Learning repository](#). The dataset is called "Online Retail" and can be found [here](#). It contains all the transactions occurring between 01/12/2010 and 09/12/2011 for a UK-based and registered online retailer.

## Load the packages

```
library(tidyverse)
library(readxl)
library(knitr)
library(ggplot2)
library(lubridate)
library(arules)
library(arulesViz)
library(plyr)
```





Open in app

Get started

```

retail <- retail %>% mutate(Description = as.factor(Description))
retail <- retail %>% mutate(Country = as.factor(Country))
retail$Date <- as.Date(retail$InvoiceDate)
retail$Time <- format(retail$InvoiceDate,"%H:%M:%S")
retail$InvoiceNo <- as.numeric(as.character(retail$InvoiceNo))

glimpse(retail)

```

```

Observations: 406,829
Variables: 10
$ InvoiceNo    <dbl> 536365, 536365, 536365, 536365, 536365, 536365, 536365, 53...
$ StockCode   <chr> "85123A", "71053", "84406B", "84029G", "84029E", "22752", ...
$ Description  <fctr> WHITE HANGING HEART T-LIGHT HOLDER, WHITE METAL LANTERN, ...
$ Quantity    <dbl> 6, 6, 8, 6, 6, 2, 6, 6, 6, 32, 6, 6, 8, 6, 6, 3, 2, 3, 3, ...
$ InvoiceDate  <dtm> 2010-12-01 08:26:00, 2010-12-01 08:26:00, 2010-12-01 08:2...
$ UnitPrice   <dbl> 2.55, 3.39, 2.75, 3.39, 3.39, 7.65, 4.25, 1.85, 1.85, 1.69...
$ CustomerID  <dbl> 17850, 17850, 17850, 17850, 17850, 17850, 17850, 17850, 17...
$ Country     <fctr> United Kingdom, United Kingdom, United Kingdom, United Ki...
$ Date        <date> 2010-12-01, 2010-12-01, 2010-12-01, 2010-12-01, 2010-12-0...
$ Time        <chr> "08:26:00", "08:26:00", "08:26:00", "08:26:00", "08:26:00"...

```

After preprocessing, the dataset includes 406,829 records and 10 fields: InvoiceNo, StockCode, Description, Quantity, InvoiceDate, UnitPrice, CustomerID, Country, Date, Time.

## What time do people often purchase online?

In order to find the answer to this question, we need to extract “hour” from the time column.

```

retail$Time <- as.factor(retail$Time)
a <- hms(as.character(retail$Time))
retail$Time = hour(a)

retail %>%
  ggplot(aes(x=Time)) +
  geom_histogram(stat="count", fill="indianred")

```





[Open in app](#)
[Get started](#)


Figure 1. Shopping time distribution

There is a clear bias between the hour of day and order volume. Most orders happened between 10:00–15:00.

### How many items each customer buy?

```
detach("package:plyr", unload=TRUE)

retail %>%
  group_by(InvoiceNo) %>%
  summarize(n_items = mean(Quantity)) %>%
  ggplot(aes(x=n_items))+
  geom_histogram(fill="indianred", bins = 100000) +
  geom_rug()+
  coord_cartesian(xlim=c(0,80))
```




[Open in app](#)[Get started](#)

Figure 2. Number of items per invoice distribution

People mostly purchased less than 10 items (less than 10 items in each invoice).

## Top 10 best sellers

```
tmp <- retail %>%
  group_by(StockCode, Description) %>%
  summarize(count = n()) %>%
  arrange(desc(count))
tmp <- head(tmp, n=10)
tmp

tmp %>%
  ggplot(aes(x=reorder(Description, count), y=count)) +
  geom_bar(stat="identity", fill="indian red") +
  coord_flip()
```



[Open in app](#)[Get started](#)

```
<chr>          <fctr> <int>
1  85123A WHITE HANGING HEART T-LIGHT HOLDER 2070
2  22423          REGENCY CAKESTAND 3 TIER 1905
3  85099B          JUMBO BAG RED RETROSPOT 1662
4  84879    ASSORTED COLOUR BIRD ORNAMENT 1418
5  47566          PARTY BUNTING 1416
6  20725          LUNCH BAG RED RETROSPOT 1358
7  22720    SET OF 3 CAKE TINS PANTRY DESIGN 1232
8    POST          POSTAGE 1196
9  20727          LUNCH BAG  BLACK SKULL. 1126
10 21212    PACK OF 72 RETROSPOT CAKE CASES 1080
>
```

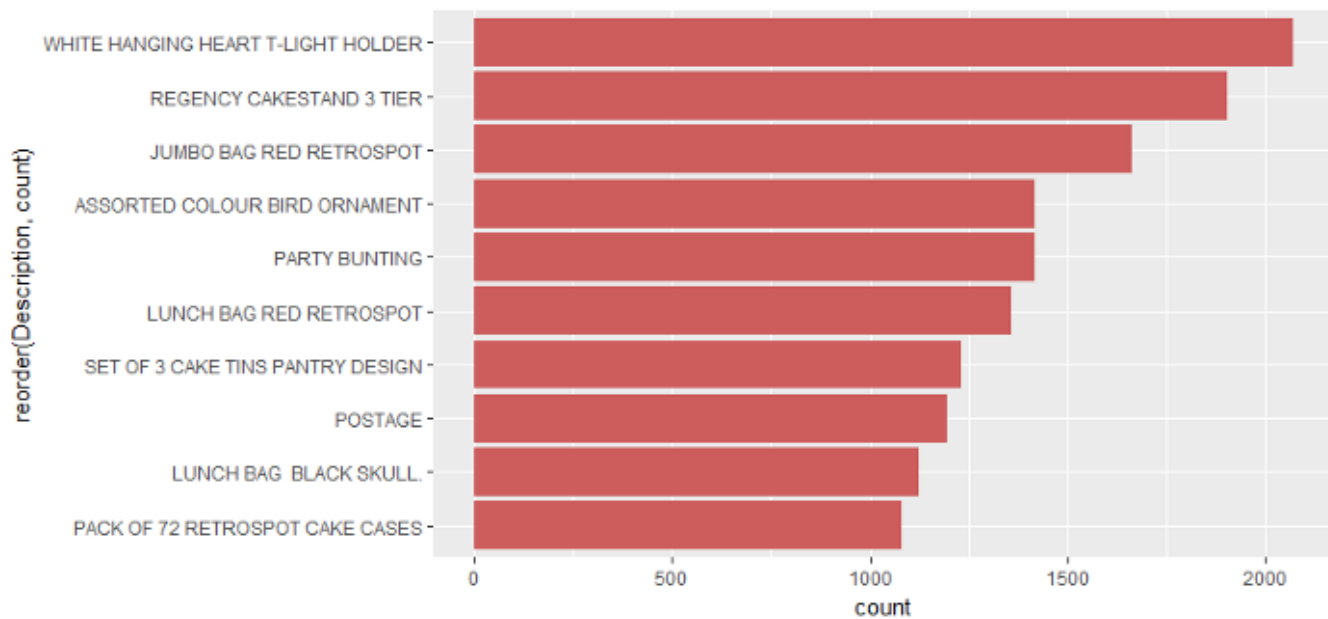


Figure 3. Top 10 best sellers

### Association rules for online retailer

Before using any rule mining algorithm, we need to transform the data from the data frame format, into transactions such that we have all the items bought together in one row. For example, this is the format we need:





Open in app

Get started

2	Touring-2000	Sport-100		
3	Mountain-200	Mountain Bottle Cage	Water Bottle	
4	Road-250	HL Road Tire	Road Tire Tube	All-Purpose E
5	Road-250	Road Bottle Cage	Water Bottle	Sport-100
6	Road-250	Road Tire Tube	HL Road Tire	Sport-100
7	Road-350-W	Long-Sleeve Logo Jersey		

Source: Microsoft

```
retail_sorted <- retail[order(retail$CustomerID),]
library(plyr)
itemList <- ddply(retail, c("CustomerID", "Date"),
                  function(df1) paste(df1$Description,
                                     collapse = ", "))
```

The function `ddply()` accepts a data frame, splits it into pieces based on one or more factors, computes on the pieces, and then returns the results as a data frame. We use “,” to separate different items.

We only need item transactions, so remove `customerID` and `Date` columns.

```
itemList$CustomerID <- NULL
itemList$Date <- NULL
colnames(itemList) <- c("items")
```

Write the data frame to a csv file and check whether our transaction format is correct.

```
write.csv(itemList, "market_basket.csv", quote = FALSE, row.names =
TRUE)
```





Open in app

Get started

5	RABBIT NI REGENCY REGENCY REGENCY REGENCY REGENCY REGENCY REGENCY REGENCY AIRLINE B AIRLINE B VICTORIA NAMASTE TRIPLE HO SMALL HE 3D DOG PI FEATHER F COAL BLA ALARM CL ALARM CLOCK BAKELIKE CH
6	SET OF 60 SET 40 HE AIRLINE B AIRLINE B AIRLINE B AIRLINE B AIRLINE B WOODLAF WOODLAF ALARM CL TRIPLE HO SINGLE AN TEA TIME 72 SWEET 60 TEATIM PACK OF 6 REGENCY REGENCY 3D DOG PI REVOLVE VINTA
7	MINI LIGH PINK GOO MADRAS AIRLINE B AIRLINE B AIRLINE B AIRLINE B BIRDCAGE CHRISTMA REGENCY REGENCY TEA TIME TEA TIME TEA TIME PINK REGI GREEN RE 3D DOG PI RABBIT NI RED TOAD TREASURE VINTA
8	CLASSIC C BICYCLE PI BOOM BO PINK NEW RED TOAD RABBIT NI WOODLAF PINK GOO CHRISTMA MINI PLAY MINI PLAYING CARDS DOLLY GIRL
9	72 SWEET 60 CAKE C 60 TEATIM 60 TEATIM PACK OF 3 PACK OF 3 PACK OF 3 PACK OF 3 SWEETIES SET OF 72 SET OF 72 60 CAKE C 60 CAKE C PACK OF 6 PACK OF 6 POSTAGE
10	PACK OF 3 PACK OF 3 MULTI HE PACK OF 3 PACK OF 3 POSTAGE
11	DOUGHNICE CREAM ICE CREAM SET OF 9 B POSTAGE
12	DOUGHNICE CREAM POSTAGE
13	PARISIEN SWEETHE PINK HEA GINGHAM RED HEAR FOOD CO LARGE HE DOORMA HANGING BROCAN PLASTERS PANTRY N RECIPE BC SET OF 3 C JAM MAKI SET OF 6 S PANTRY C DOORMA 16 PIECE C SMALL WH BLAC
14	CHOCOLA METAL SIC RETRO MC RETRO PU TEA BAG F PINK PUR PLASTERS PLASTERS CHOCOLA RED HARM 4 TRADITH BATHROO POSTAGE UNION JA UNION JA BLUE POLI BLUE POLKADOT PASSPORT COVER
15	WOODEN PINK DOU STRAWBEI CERAMIC WOODEN REGENCY DELUXE SE WELCOME LOVE BUIL BATH BUIL HOME BUI CAT BOWI BIG DOUG DOLLY GIR LIGHT GARLAND BUTTERFILES PINK
16	POSTAGE DELUXE SE PINK HEA BAKING SI VINTAGE Manual Manual Manual Manual Manual Manual
17	CERAMIC CERAMIC BLUE HARI PINK DOG PINK HEA LANTERN METAL SIGN TAKE IT OR LEAVE IT
18	PINK HEA CERAMIC LANTERN PINK DOG CERAMIC METAL SIC BLUE HARI POSTAGE PINK HEA CERAMIC LANTERN PINK DOG CERAMIC METAL SIC BLUE HARMONICA IN BOX
19	ANTIQUE PANTRY N PANTRY S SET OF 3 R SMALL GL DOORMA SET OF 4 P PANTRY W BAKING SI IVORY KIT SET OF 3 C REGENCY WOODEN LIGHT GAF FAIRY CAK SPOTTY BI SET OF 4 E POSTAGE
20	OPEN CLOSET OF 6 S SET OF 3 R RED TOAD CHILDS BR CHILDS BR SET OF 3 C SET OF TE LOVE BUIL HOLIDAY F BATH BUIL WOODEN LIGHT GAF POSTAGE
21	PETIT TRA PANTRY R PANTRY P WOODLAF PINK RAR MINT KITC SET 12 COLZINC HEA VICTORIA IVORY KIT GLASS BO BLUE STRU SET 12 COL CHILDS BR POSTAGE

Perfect! Now we have our transaction dataset, and it shows the matrix of items being bought together. We don't actually see how often they are bought together, and we don't see rules either. But we are going to find out.

Let's have a closer look at how many transactions we have and what they are.

```
tr <- read.transactions('market_basket.csv', format = 'basket',
  sep=',')
tr
summary(tr)
```







Open in app

Get started

```
transactions as itemMatrix in sparse format with
19296 rows (elements/itemsets/transactions) and
7881 columns (items) and a density of 0.002200461
```

```
most frequent items:
```

```
WHITE HANGING HEART T-LIGHT HOLDER      REGENCY CAKESTAND 3 TIER
                                1772                                1667
      JUMBO BAG RED RETROSPOT              PARTY BUNTING
                                1445                                1279
ASSORTED COLOUR BIRD ORNAMENT              (other)
                                1239                                327226
```

```
element (itemset/transaction) length distribution:
sizes
```

1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
2247	1177	848	762	724	660	614	595	584	553	574	507	490	507	503	504
17	18	19	20	21	22	23	24	25	26	27	28	29	30	31	32
452	415	474	420	383	309	311	271	236	253	223	204	226	218	174	146
33	34	35	36	37	38	39	40	41	42	43	44	45	46	47	48
139	145	130	112	116	88	104	94	91	86	94	60	68	74	68	65
49	50	51	52	53	54	55	56	57	58	59	60	61	62	63	64
52	50	60	51	41	53	51	36	23	40	37	30	31	23	22	24
65	66	67	68	69	70	71	72	73	74	75	76	77	78	79	80
17	27	32	22	17	25	17	20	18	12	13	19	14	7	9	18
81	82	83	84	85	86	87	88	89	90	91	92	93	94	95	96
17	11	10	8	12	10	15	7	7	9	6	7	8	5	4	5
97	98	99	100	101	102	103	104	105	106	107	108	109	110	111	112
5	3	3	3	5	5	5	2	3	3	8	5	6	3	3	1
113	114	115	116	117	118	119	120	121	122	123	125	126	127	131	132
2	2	1	4	6	3	1	2	1	3	3	4	2	1	1	1

We see 19,296 transactions, and this is the number of rows as well. There are 7,881 items — remember items are the product descriptions in our original dataset. Transactions here are the collections or subsets of these 7,881 items.

The summary gives us some useful information:

- density: The percentage of non-empty cells in the sparse matrix. In another words, the total number of items that are purchased divided by the total number of possible items in that matrix. We can calculate how many items were purchased using density like so:  
 **$19296 \times 7881 \times 0.0022$**
- The most frequent items should be the same as our results in Figure 3.
- Looking at the size of the transactions: 2247 transactions were for just 1 item, 1147 transactions for 2 items, all the way up to the biggest transaction: 1 transaction for 420 items. This indicates that most customers buy a small number of items in each





Open in app

Get started

```
includes extended item information - examples:
      labels
1          1 HANGER
2      10 COLOUR SPACEBOY PEN
3     12 COLOURED PARTY BALLOONS
> |
```

- The distribution of the data is right skewed.

Let's have a look at the item frequency plot, which should be in aligned with Figure 3.

```
itemFrequencyPlot(tr, topN=20, type='absolute')
```




Figure 4. A bar plot of the support of the 20 most frequent items bought.

### Create some rules

- We use the Apriori algorithm in Arules library to mine frequent itemsets and association rules. The algorithm employs level-wise search for frequent itemsets.
- We pass  $\text{supp}=0.001$  and  $\text{conf}=0.8$  to return all the rules that have a support of at least 0.1% and confidence of at least 80%.





Open in app

Get started

```
rules <- apriori(tr, parameter = list(supp=0.001, conf=0.8))
rules <- sort(rules, by='confidence', decreasing = TRUE)
summary(rules)
```

```
> summary(rules)
set of 89697 rules

rule length distribution (lhs + rhs):sizes
  2    3    4    5    6    7    8    9   10
103 3206 9909 26451 31144 14599 3464  700  121

  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
 2.000  5.000   6.000   5.641  6.000  10.000

summary of quality measures:
  support      confidence      lift      count
Min.   :0.001036   Min.   :0.8000   Min.   : 8.711   Min.   : 20.00
1st Qu.:0.001088   1st Qu.:0.8333   1st Qu.:19.052   1st Qu.: 21.00
Median :0.001192   Median :0.8750   Median :24.495   Median : 23.00
Mean   :0.001382   Mean   :0.8827   Mean   :49.558   Mean   : 26.67
3rd Qu.:0.001503   3rd Qu.:0.9231   3rd Qu.:42.265   3rd Qu.: 29.00
Max.   :0.018242   Max.   :1.0000   Max.   :622.452   Max.   :352.00

mining info:
 data ntransactions support confidence
  tr          19296    0.001         0.8
> |
```

The summary of the rules gives us some very interesting information:

- The number of rules: 89,697.
- The distribution of rules by length: a length of 6 items has the most rules.
- The summary of quality measures: ranges of support, confidence, and lift.
- The information on data mining: total data mined, and the minimum parameters we set earlier.

We have 89,697 rules. I don't want to print them all, so let's inspect the top 10.

```
inspect(rules[1:10])
```



[Open in app](#)[Get started](#)

```

[2] {WOBBLY CHICKEN} => {METAL} 0.001431078 1
[3] {DECOUPAGE}      => {GREETING CARD} 0.001191957 1
[4] {BILLBOARD FONTS DESIGN} => {WRAP} 0.001502902 1
[5] {WOBBLY RABBIT}  => {DECORATION} 0.001762023 1
[6] {WOBBLY RABBIT}  => {METAL} 0.001762023 1
[7] {BLACK TEA}      => {SUGAR JARS} 0.002332090 1
[8] {BLACK TEA}      => {COFFEE} 0.002332090 1
[9] {CHOCOLATE SPOTS} => {SWISS ROLL TOWEL} 0.002176617 1
[10] {ART LIGHTS}     => {FUNK MONKEY} 0.002021144 1
lift      count
[1] 385.92000 28
[2] 385.92000 28
[3] 344.57143 23
[4] 622.45161 29
[5] 385.92000 34
[6] 385.92000 34
[7] 212.04396 45
[8]  61.06329 45
[9] 410.55319 42
[10] 494.76923 39
> |

```

The interpretation is pretty straight forward:

- 100% customers who bought “WOBBLY CHICKEN” also bought “DECORATION”.
- 100% customers who bought “BLACK TEA” also bought “SUGAR JAR”.

And plot these top 10 rules.

```

topRules <- rules[1:10]
plot(topRules)

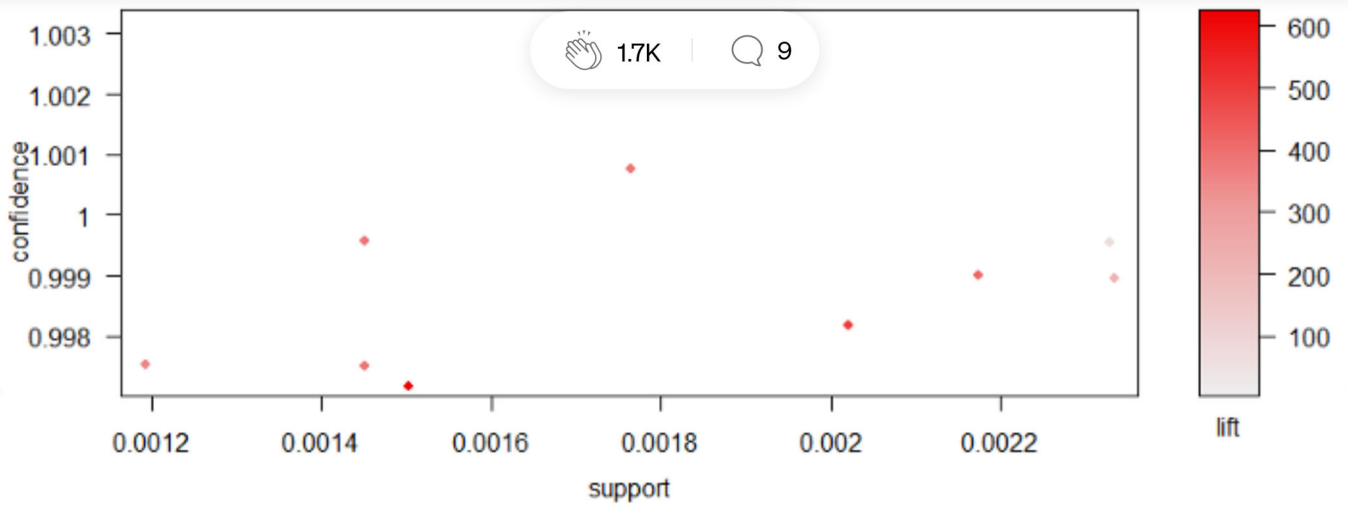
```





Open in app

Get started



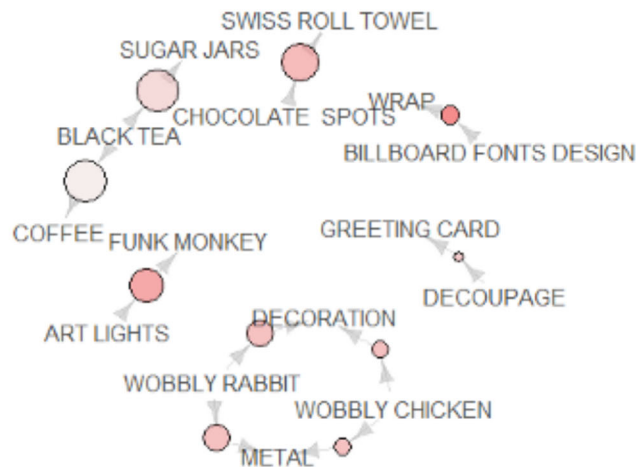
by Towards Data Science

Every Thursday, the Variable delivers the very best of Towards Data Science: from hands-on tutorials and cutting-edge research to original features you don't want to miss. [Take a look.](#)

Your email

### Graph for 10 rules

size: support (0.001 - 0.002)  
color: lift (61.063 - 622.452)



```
plot(topRules, method = "grouped")
```



[Open in app](#)[Get started](#)

## Summary

In this post, we have learned how to perform Market Basket Analysis in R and how to interpret the results. If you want to implement them in Python, [Mlxtend](#) is a Python library that has an implementation of the Apriori algorithm for this sort of application. You can find an introductory tutorial [here](#).

If you would like the R Markdown file used to make this blog post, you can find [here](#).

reference: [R and Data Mining](#)

